

# Telecom Churn Analysis

```
In [1]: #import the required libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.ticker as mtick
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
```

## Load the data file

```
In [2]: telco_base_data = pd.read_csv('Telco-Customer-Churn.csv')
```

Look at the top 5 records of data

```
In [3]: telco_base_data.columns
```

```
Out[3]: Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
              'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
              'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
              'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
              'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
              dtype='object')
```

```
In [4]: telco_base_data.head()
```

```
Out[4]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	Mul
--	------------	--------	---------------	---------	------------	--------	--------------	-----

0	7590-VHVEG	Female	0	Yes	No	1	No	
---	------------	--------	---	-----	----	---	----	--


1	5575-GNVDE	Male	0	No	No	34	Yes	
---	------------	------	---	----	----	----	-----	--

2	3668-QPYBK	Male	0	No	No	2	Yes	
---	------------	------	---	----	----	---	-----	--

3	7795-CFOCW	Male	0	No	No	45	No	
---	------------	------	---	----	----	----	----	--

4	9237-HQITU	Female	0	No	No	2	Yes	
---	------------	--------	---	----	----	---	-----	--

5 rows × 21 columns



Check the various attributes of data like shape (rows and cols), Columns, datatypes

```
In [5]: telco_base_data.shape
```

Out[5]: (7043, 21)

```
In [6]: telco_base_data.columns.values
```

```
Out[6]: array(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',  
             'tenure', 'PhoneService', 'MultipleLines', 'InternetService',  
             'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',  
             'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',  
             'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges',  
             'TotalCharges', 'Churn'], dtype=object)
```

```
In [7]: # Checking the data types of all the columns  
telco_base_data.dtypes
```

```
Out[7]: customerID      object  
gender      object  
SeniorCitizen  int64  
Partner      object  
Dependents    object  
tenure      int64  
PhoneService  object  
MultipleLines object  
InternetService object  
OnlineSecurity object  
OnlineBackup  object  
DeviceProtection object  
TechSupport   object  
StreamingTV   object  
StreamingMovies object  
Contract      object  
PaperlessBilling object  
PaymentMethod object  
MonthlyCharges float64  
TotalCharges  object  
Churn         object  
dtype: object
```

```
In [8]: # Check the descriptive statistics of numeric variables  
telco_base_data.describe()
```

```
Out[8]:
```


	SeniorCitizen	tenure	MonthlyCharges
<b>count</b>	7043.000000	7043.000000	7043.000000
<b>mean</b>	0.162147	32.371149	64.761692
<b>std</b>	0.368612	24.559481	30.090047
<b>min</b>	0.000000	0.000000	18.250000
<b>25%</b>	0.000000	9.000000	35.500000
<b>50%</b>	0.000000	29.000000	70.350000
<b>75%</b>	0.000000	55.000000	89.850000
<b>max</b>	1.000000	72.000000	118.750000

SeniorCitizen is actually a categorical hence the 25%-50%-75% distribution is not proper

75% customers have tenure less than 55 months

Average Monthly charges are USD 64.76 whereas 25% customers pay more than USD 89.85 per month

```
In [9]: telco_base_data['Churn'].value_counts().plot(kind='barh', figsize=(8, 6))
plt.xlabel("Count", labelpad=14)
plt.ylabel("Target Variable", labelpad=14)
plt.title("Count of TARGET Variable per category", y=1.02);
```



```
In [10]: 100*telco_base_data['Churn'].value_counts()/len(telco_base_data['Churn'])
```

```
Out[10]: Churn
No      73.463013
Yes     26.536987
Name: count, dtype: float64
```

```
In [11]: telco_base_data['Churn'].value_counts()
```

```
Out[11]: Churn
No      5174
Yes     1869
Name: count, dtype: int64
```

- Data is highly imbalanced, ratio = 73:27
- So we analyse the data with other features while taking the target values separately to get some insights.

```
In [12]: # Concise Summary of the dataframe, as we have too many columns, we are using th
telco_base_data.info(verbose = True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines           7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity         7043 non-null   object
10  OnlineBackup           7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7043 non-null   object
20  Churn                  7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
In [13]: missing = pd.DataFrame(((telco_base_data.isnull().sum())*100/telco_base_data.shap
```


```
In [14]: missing
```

Out[14]:

	index	0
0	customerID	0.0
1	gender	0.0
2	SeniorCitizen	0.0
3	Partner	0.0
4	Dependents	0.0
5	tenure	0.0
6	PhoneService	0.0
7	MultipleLines	0.0
8	InternetService	0.0
9	OnlineSecurity	0.0
10	OnlineBackup	0.0
11	DeviceProtection	0.0
12	TechSupport	0.0
13	StreamingTV	0.0
14	StreamingMovies	0.0
15	Contract	0.0
16	PaperlessBilling	0.0
17	PaymentMethod	0.0
18	MonthlyCharges	0.0
19	TotalCharges	0.0
20	Churn	0.0

In [15]: `import matplotlib.pyplot as plt`  
`import seaborn as sns`

```
plt.figure(figsize=(16, 5))
ax = sns.pointplot(x='index', y=0, data=missing)
plt.xticks(rotation=90, fontsize=7)
plt.title("Percentage of Missing values")
plt.ylabel("PERCENTAGE")
plt.show()
```



## Missing Data - Initial Intuition

- Here, we don't have any missing data.

General Thumb Rules:

- For features with less missing values- can use regression to predict the missing values or fill with the mean of the values present, depending on the feature.
- For features with very high number of missing values- it is better to drop those columns as they give very less insight on analysis.
- As there's no thumb rule on what criteria do we delete the columns with high number of missing values, but generally you can delete the columns, if you have more than 30-40% of missing values. But again there's a catch here, for example, Is\_Car & Car\_Type, People having no cars, will obviously have Car\_Type as NaN (null), but that doesn't make this column useless, so decisions has to be taken wisely.

## Data Cleaning

1. Create a copy of base data for manipulation & processing

```
In [16]: telco_data = telco_base_data.copy()
```

2. Total Charges should be numeric amount. Let's convert it to numerical data type

The errors='coerce' parameter is used to handle errors.

When errors is set to 'coerce', it means that if there are any errors encountered while converting data to numeric, those errors will be replaced with NaN (Not a Number) values.

```
In [17]: telco_data.TotalCharges = pd.to_numeric(telco_data.TotalCharges, errors='coerce')
telco_data.isnull().sum()
```

```
Out[17]: customerID      0
gender      0
SeniorCitizen  0
Partner      0
Dependents    0
tenure      0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport    0
StreamingTV    0
StreamingMovies  0
Contract      0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges   11
Churn         0
dtype: int64
```


**3.** As we can see there are 11 missing values in TotalCharges column. Let's check these records

```
In [18]: telco_data.loc[telco_data ['TotalCharges'].isnull() == True]
```

Out[18]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService
488	4472-LVYGI	Female	0	Yes	Yes	0	No
753	3115-CZMZD	Male	0	No	Yes	0	Yes
936	5709-LVOEQ	Female	0	Yes	Yes	0	Yes
1082	4367-NUYAO	Male	0	Yes	Yes	0	Yes
1340	1371-DWPAZ	Female	0	Yes	Yes	0	No
3331	7644-OMVMY	Male	0	Yes	Yes	0	Yes
3826	3213-VVOLG	Male	0	Yes	Yes	0	Yes
4380	2520-SGTTA	Female	0	Yes	Yes	0	Yes
5218	2923-ARZLG	Male	0	Yes	Yes	0	Yes
6670	4075-WKNIU	Female	0	Yes	Yes	0	Yes
6754	2775-SEFEE	Male	0	No	Yes	0	Yes

11 rows × 21 columns



#### 4. Missing Value Treatement

Since the % of these records compared to total dataset is very low ie 0.15%, it is safe to ignore them from further processing.

```
In [19]: #Removing missing values
telco_data.dropna(how = 'any', inplace = True)

#telco_data.fillna(0)
```

5. Divide customers into bins based on tenure e.g. for tenure < 12 months: assign a tenure group if 1-12, for tenure between 1 to 2 Yrs, tenure group of 13-24; so on...

```
In [20]: # Get the max tenure
print(telco_data['tenure'].max()) #72
```

72

```
In [21]: # Group the tenure in bins of 12 months
labels = ["{0} - {1}".format(i, i + 11) for i in range(1, 72, 12)]
```



```
telco_data['tenure_group'] = pd.cut(telco_data.tenure, range(1, 80, 12), right=False)
```

by cutting the 'tenure' column into intervals defined by the range(1, 80, 12) function, which generates intervals starting from 1 and ending at 72 (since the range stops before the end value). The right=False parameter means that the intervals are closed on the left and open on the right.

```
In [22]: telco_data['tenure_group'].value_counts()
```

```
Out[22]: tenure_group
1 - 12      2175
61 - 72     1407
13 - 24     1024
25 - 36      832
49 - 60      832
37 - 48      762
Name: count, dtype: int64
```

## 6. Remove columns not required for processing

```
In [23]: #drop column customerID and tenure
telco_data.drop(columns=['customerID', 'tenure'], axis=1, inplace=True)
telco_data.head()
```

```
Out[23]:
```

	gender	SeniorCitizen	Partner	Dependents	PhoneService	MultipleLines	InternetService
--	--------	---------------	---------	------------	--------------	---------------	-----------------


0	Female	0	Yes	No	No	No phone service	
---	--------	---	-----	----	----	------------------	--

1	Male	0	No	No	Yes	No	
---	------	---	----	----	-----	----	--

2	Male	0	No	No	Yes	No	
---	------	---	----	----	-----	----	--

3	Male	0	No	No	No	No phone service	
---	------	---	----	----	----	------------------	--

4	Female	0	No	No	Yes	No	Fiber
---	--------	---	----	----	-----	----	-------



# Data Exploration



## 1. Plot distribution of individual predictors by churn



## Univariate Analysis



```
In [24]: warnings.simplefilter(action='ignore', category=FutureWarning)
```



```
In [25]: for i, predictor in enumerate(telco_data.drop(columns=['Churn', 'TotalCharges'],
plt.figure(i)
```



```
sns.countplot(data=telco_data, x=predictor, hue='Churn')
```












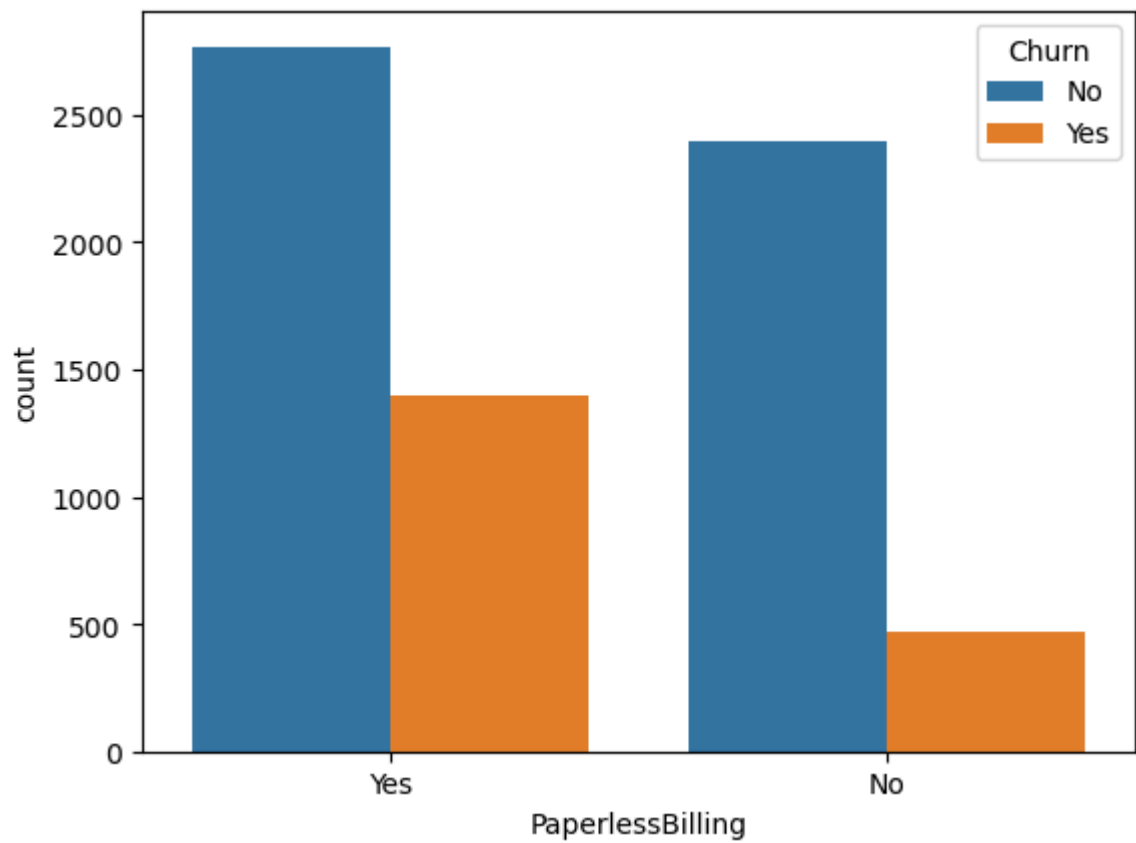















2. Convert the target variable 'Churn' in a binary numeric variable i.e. Yes=1 ; No = 0

```
In [26]: telco_data['Churn'] = np.where(telco_data.Churn == 'Yes',1,0)
```

```
In [27]: telco_data.head()
```

```
Out[27]:
```

	gender	SeniorCitizen	Partner	Dependents	PhoneService	MultipleLines	InternetSe
--	--------	---------------	---------	------------	--------------	---------------	------------


0	Female	0	Yes	No	No	No phone service	
---	--------	---	-----	----	----	------------------	--

1	Male	0	No	No	Yes	No	
---	------	---	----	----	-----	----	--

2	Male	0	No	No	Yes	No	
---	------	---	----	----	-----	----	--

3	Male	0	No	No	No	No phone service	
---	------	---	----	----	----	------------------	--

4	Female	0	No	No	Yes	No	Fiber
---	--------	---	----	----	-----	----	-------




3. Convert all the categorical variables into dummy variables

```
In [28]: telco_data_dummies = pd.get_dummies(telco_data)
telco_data_dummies.head()
```

Out[28]:

	SeniorCitizen	MonthlyCharges	TotalCharges	Churn	gender_Female	gender_Male	I
0	0	29.85	29.85	0	True	False	
1	0	56.95	1889.50	0	False	True	
2	0	53.85	108.15	1	False	True	
3	0	42.30	1840.75	0	False	True	
4	0	70.70	151.65	1	True	False	


5 rows × 51 columns



## 9. Relationship between Monthly Charges and Total Charges

```
In [29]: sns.lmplot(data=telco_data_dummies, x='MonthlyCharges', y='TotalCharges', fit_re
```

Out[29]: <seaborn.axisgrid.FacetGrid at 0x229de97dc90>



Total Charges increase as Monthly Charges increase - as expected.

## 10. Churn by Monthly Charges and Total Charges


```
In [30]: Mth = sns.kdeplot(telco_data_dummies.MonthlyCharges[(telco_data_dummies["Churn"]  
color="Red", shade = True)  
Mth = sns.kdeplot(telco_data_dummies.MonthlyCharges[(telco_data_dummies["Churn"]
```

```

        ax = Mth, color="Blue", shade= True)
Mth.legend(["No Churn", "Churn"], loc='upper right')
Mth.set_ylabel('Density')
Mth.set_xlabel('Monthly Charges')
Mth.set_title('Monthly charges by churn')

```

Out[30]: Text(0.5, 1.0, 'Monthly charges by churn')




**Insight:** Churn is high when Monthly Charges are high

```

In [31]: Tot = sns.kdeplot(telco_data_dummies.TotalCharges[(telco_data_dummies["Churn"] =
        color="Red", shade = True)
Tot = sns.kdeplot(telco_data_dummies.TotalCharges[(telco_data_dummies["Churn"] =
        ax = Tot, color="Blue", shade= True)
Tot.legend(["No Churn", "Churn"], loc='upper right')
Tot.set_ylabel('Density')
Tot.set_xlabel('Total Charges')
Tot.set_title('Total charges by churn')

```

Out[31]: Text(0.5, 1.0, 'Total charges by churn')




**Surprising insight** as higher Churn at lower Total Charges

However if we combine the insights of 3 parameters i.e. Tenure, Monthly Charges & Total Charges then the picture is bit clear :- Higher Monthly Charge at lower tenure results into lower Total Charge. Hence, all these 3 factors viz **Higher Monthly Charge**, **Lower tenure** and **Lower Total Charge** are linkd to **High Churn**.

## 11. Build a corelation of all predictors with 'Churn'

```
In [32]: plt.figure(figsize=(20,8))
telco_data_dummies.corr()['Churn'].sort_values(ascending = False).plot(kind='bar')
```

Out[32]: <Axes: >



### Derived Insight:

**HIGH** Churn seen in case of **Month to month contracts, No online security, No Tech support, First year of subscription** and **Fibre Optics Internet**


**LOW** Churn is seen in case of **Long term contracts, Subscriptions without internet service** and **The customers engaged for 5+ years**

Factors like **Gender, Availability of PhoneService** and **# of multiple lines** have almost **NO** impact on Churn

This is also evident from the **Heatmap** below

```
In [33]: plt.figure(figsize=(12,12))
sns.heatmap(telco_data_dummies.corr(), cmap="Paired")
```

```
Out[33]: <Axes: >
```



## Bivariate Analysis

```
In [34]: new_df1_target0=telco_data.loc[telco_data["Churn"]==0]
new_df1_target1=telco_data.loc[telco_data["Churn"]==1]
```

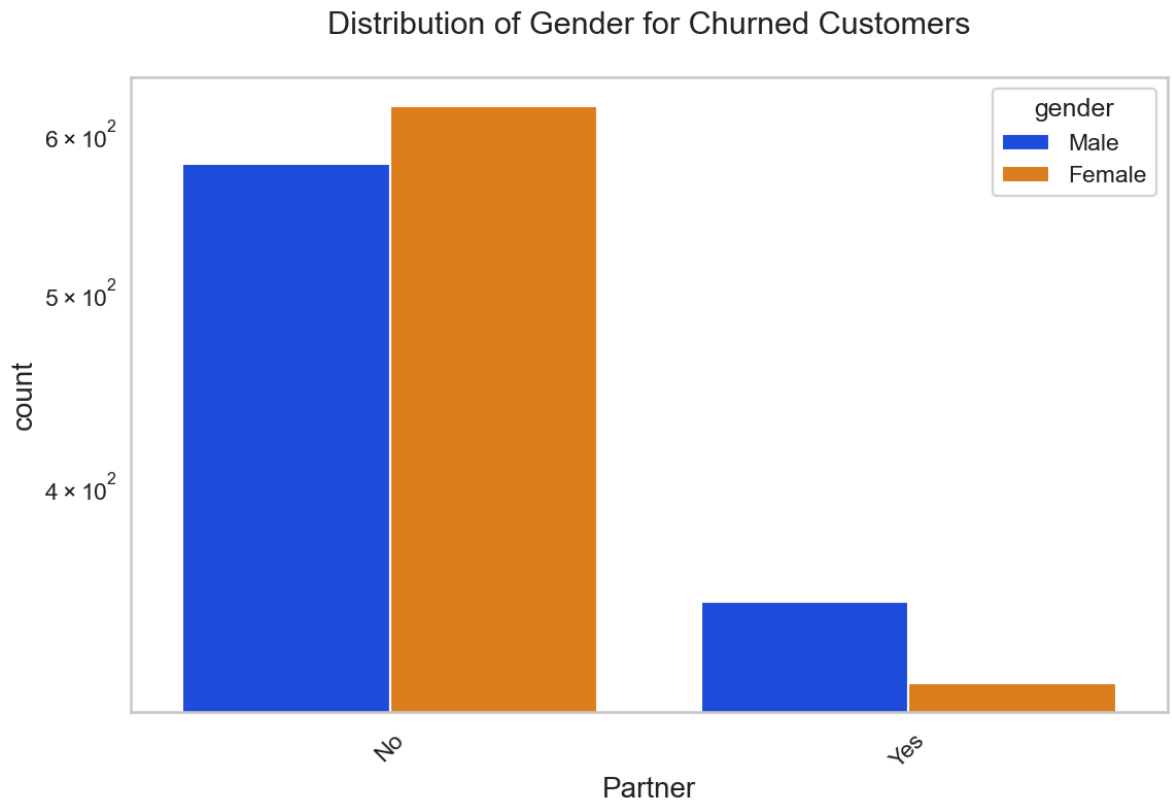
```
In [35]: def uniplot(df,col,title,hue =None):

    sns.set_style('whitegrid')
    sns.set_context('talk')
    plt.rcParams["axes.labelsize"] = 20
    plt.rcParams['axes.titlesize'] = 22
    plt.rcParams['axes.titlepad'] = 30 #These lines set the size and padding of

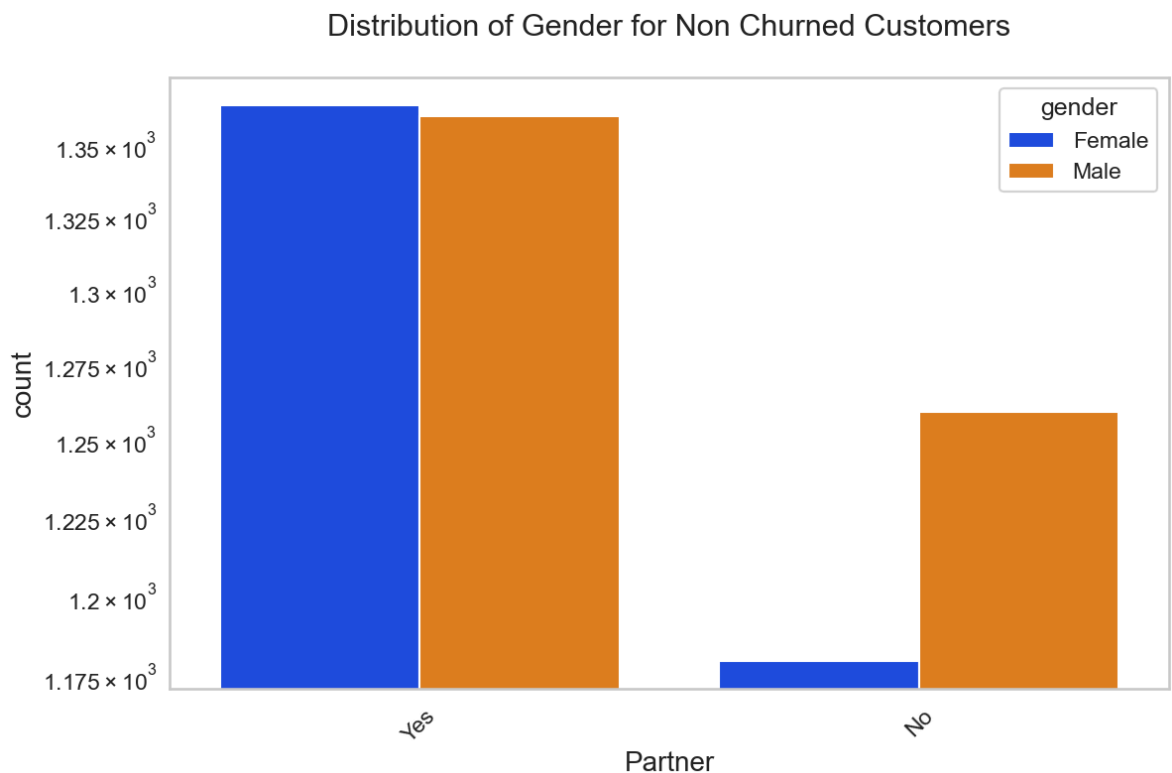
    temp = pd.Series(data = hue)
    fig, ax = plt.subplots()
    width = len(df[col].unique()) + 7 + 4*len(temp.unique())
    #This line calculates the width of the plot based on the number of unique va
    #in the specified column (col) and the unique values in the hue parameter.
    fig.set_size_inches(width , 8)
    plt.xticks(rotation=45)
    plt.yscale('log')
```

```
plt.title(title)
ax = sns.countplot(data = df, x= col, order=df[col].value_counts().index,hue
#This line creates the count plot using seaborn's countplot function. It spe
#the column (col) to plot on the x-axis, the order of the categories based o
#the hue parameter for further categorization, and the color palette ('brigh
plt.show()
```

In [36]: `#The uniplot() function you've provided is intended to create a count plot using uniplot(new_df1_target1,col='Partner',title='Distribution of Gender for Churned`




In [37]: `uniplot(new_df1_target0,col='Partner',title='Distribution of Gender for Non Churned`






```
In [38]: uniplot(new_df1_target1,col='PaymentMethod',title='Distribution of PaymentMethod
```




```
In [39]: uniplot(new_df1_target1,col='Contract',title='Distribution of Contract for Churn
```




```
In [40]: uniplot(new_df1_target1,col='TechSupport',title='Distribution of TechSupport for
```

Distribution of TechSupport for Churned Customers



```
In [41]: uniplot(new_df1_target1,col='SeniorCitizen',title='Distribution of SeniorCitizen')
```


Distribution of SeniorCitizen for Churned Customers



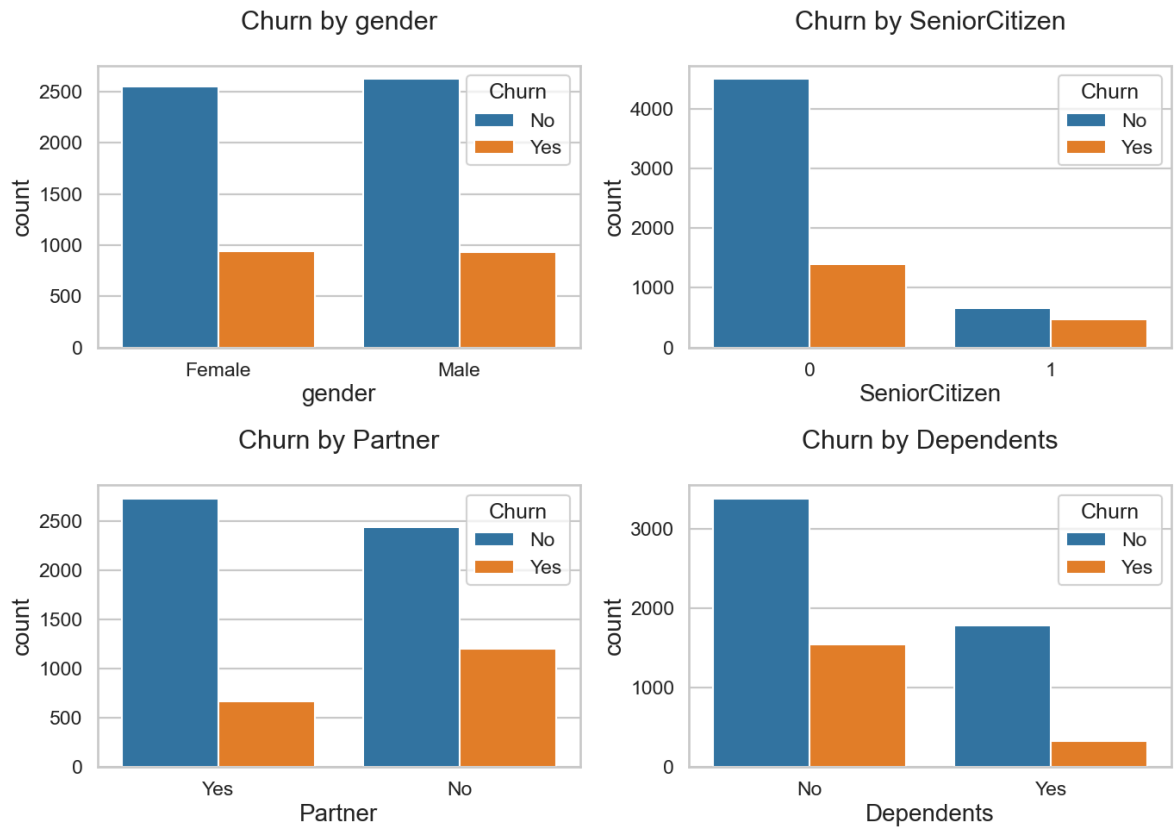
```
In [42]: # 1. Churn Rate Analysis
churn_rate = telco_base_data['Churn'].value_counts(normalize=True) * 100
plt.figure(figsize=(6, 6))
churn_rate.plot(kind='pie', autopct='%1.1f%%', labels=['Retained', 'Churned'])
plt.title('Churn Rate')
```

```
plt.ylabel('')  
plt.show()
```


## Churn Rate



```
In [43]: # 2. Demographic Analysis  
demographics = ['gender', 'SeniorCitizen', 'Partner', 'Dependents']  
plt.figure(figsize=(14, 10))  
for i, column in enumerate(demographics, 1):  
    plt.subplot(2, 2, i)  
    sns.countplot(data=telco_base_data, x=column, hue='Churn')  
    plt.title(f'Churn by {column}')  
plt.tight_layout()  
plt.show()
```

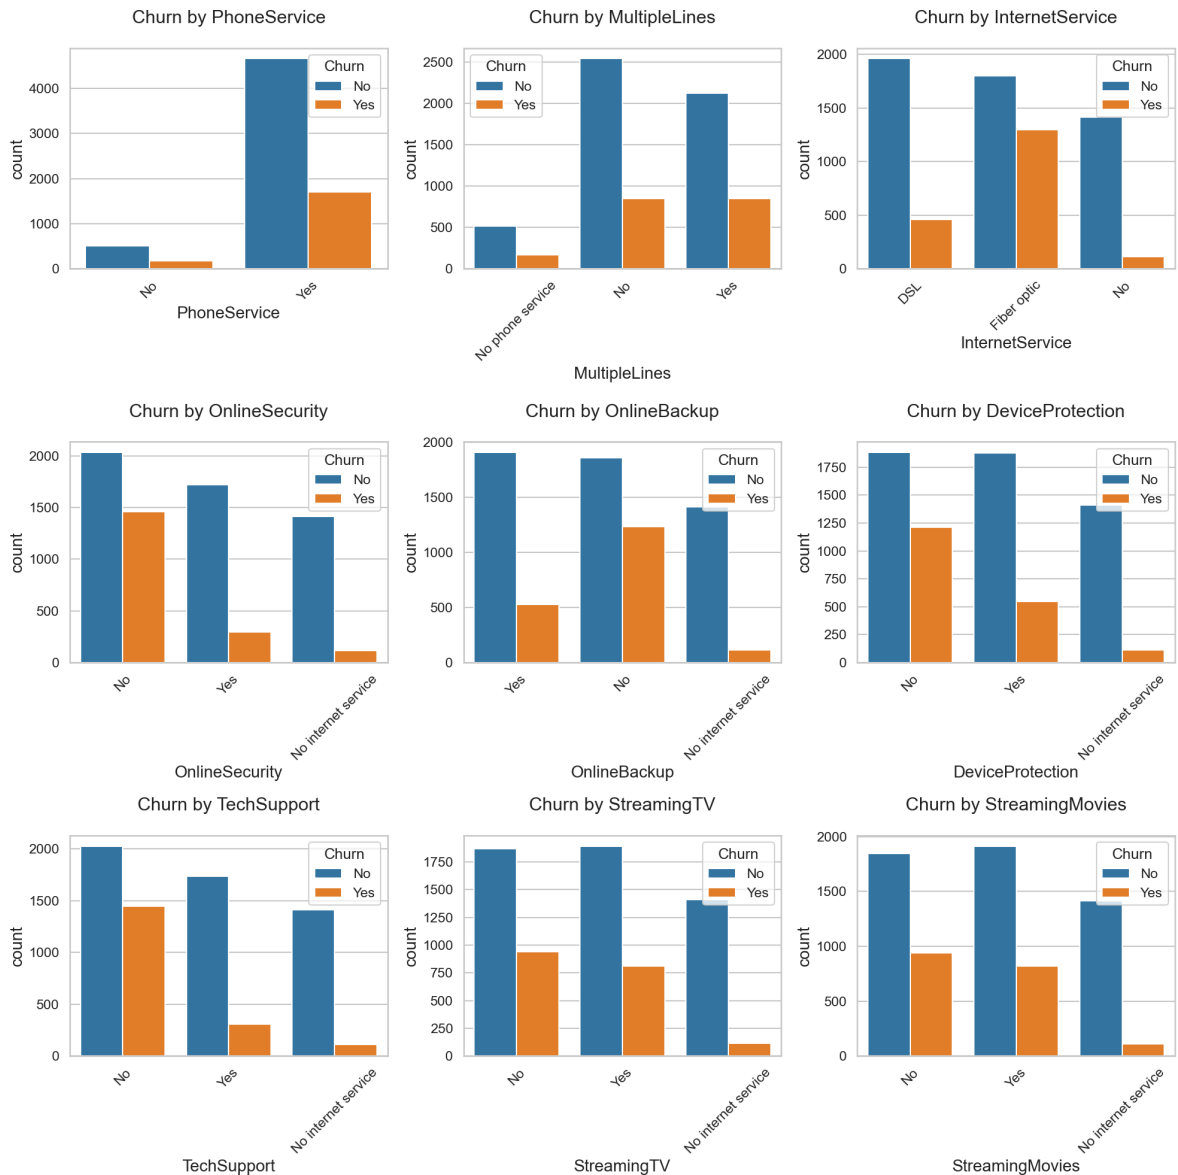


```
In [44]: # 3. Tenure Analysis
plt.figure(figsize=(10, 6))
sns.boxplot(data=telco_base_data, x='Churn', y='tenure')
plt.title('Tenure Distribution by Churn')
plt.show()
```





```
In [45]: # 4. Service Usage Analysis
services = ['PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
            'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'S
```

```
plt.figure(figsize=(20, 20))
for i, service in enumerate(services, 1):
    plt.subplot(3, 3, i)
    sns.countplot(data=telco_base_data, x=service, hue='Churn')
    plt.title(f'Churn by {service}')
    plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
In [46]: # 5. Contract and Billing Analysis
plt.figure(figsize=(14, 6))
sns.countplot(data=telco_base_data, x='Contract', hue='Churn')
plt.title('Churn by Contract Type')
plt.xticks(rotation=45)
plt.show()

plt.figure(figsize=(12, 6))
sns.boxplot(data=telco_base_data, x='Churn', y='MonthlyCharges')
plt.title('Monthly Charges Distribution by Churn')
plt.show()
```



## CONCLUSION

1. Electronic check medium are the highest churners
2. Contract Type - Monthly customers are more likely to churn because of no contract terms, as they are free to go customers.
3. No Online security, No Tech Support category are high churners
4. Non senior Citizens are high churners

```
In [47]: telco_data_dummies.to_csv('tel_churn.csv')
```

```
In [55]: import pandas as pd
import sqlalchemy as sa
```

```
In [58]: from sqlalchemy import create_engine

# Define the database connection URL
# Replace 'username', 'password', 'host', and 'database_name' with your MySQL se
db_url = 'mysql+pymysql://root:2002@localhost/Customer_churn'

# Create an engine
engine = create_engine(db_url)

# Test the connection
try:
    engine.connect()
    print("Connection successful!")
except Exception as e:
    print("Connection failed:", e)
```

Connection successful!

```
In [63]: # Write DataFrame to SQL table
try:
    telco_base_data.to_sql(name='telco_customer_churn', con=engine, if_exists='f
    print("DataFrame successfully written to SQL table.")
except Exception as e:
    print("DataFrame to SQL table failed:", e)
```

DataFrame successfully written to SQL table.

```
In [ ]:
```